

Performance Optimization


Justin Anguiano

August 20, 2019

Goal — Improve Physics analysis performance

Use different C/Python approaches to obtain a low programming overhead but high performance

Today: a look at multithreading and multiprocessing in C++ ROOT



Multithreading

Multiprocessing-

- adds CPUs to increase performance
- processes run concurrently
- creating/managing processes is resource intensive
- processes do not share memory

Multithreading-

- adds threads to a single CPU to increase performance
- multithreading is economic for the system
- threads share memory

There are parallelization tools to exploit both of these approaches built into root!

Definitions:

thread of execution is the smallest sequence of programmed instructions that can be managed independently by a scheduler

The number of threads possible depend on the number of cores in the CPU

process is the instance of a computer program that is being executed by one or many threads

Test 1: Local multi-threading

Task to be optimized: Read File(s) containing a TTree, produce histograms from elements of the tree(s), write histograms to a TFile

Three histograms are produced: TH1D: track p_T weighted by $1/p_T$, TH1D: track p_z , TH2D, track p_x vs p_y

First attempt using ROOT6 parallelization on local machine (laptop) methods:

- Sequential run over a test file with TTreeReader (Compiled and Interpreted)
- Parallelized run over a test file with TTreeReader (Compiled and Interpreted)


Test File Details

- code adapted from example: https://root.cern.ch/doc/v612/imt101__parTreeProcessing_8C.html
- uses fake event file containing 48,000 events; each with some number of tracks per event
- total tracks overall $\approx 2.4e + 7$

Test 1: Local multi-threading

First Test Results:

- Time is Mean time \pm stdev over 10 trials
- no guarantee of system releasing requested resources (nthreads) so perform multiple trials
- data point at nThreads = 0 is the basic sequential program
- high threadcount performance possibly bottlenecked by system?
- Takeaway: Compiling is important!
 \approx Factor of 2 improvement



Test 2: UNL multi-threading

Second attempt using ROOT6 parallelization versus classic(Make Class) sequential program (my current approach) on t3.unl.edu

Methods:

- Sequential run over multiple test files with TTreeReader (Compiled Only)
- Parallelized run over multiple test files with TTreeReader (Compiled Only)
- Sequential run over multiple test files with MakeClass (Compiled and Interpreted)

Test File set Details:

- using 9 files from Dataset:
/SingleMuon/Run2018D-PromptReco-v2/A0D
 - Physics group: NoGroup Creation time: 2018-08-01 13:16:41 Status: VALID Type: data Dataset size: 150070502441346 (150.1TB) Number of blocks: 267 Number of events: 511823047 Number of files: 45330
- 1,576,737 events each with at least 1 conversion per event
- total tracks overall (exactly 2 per conversion) = 14,190,956

Test 2: UNL multi-threading

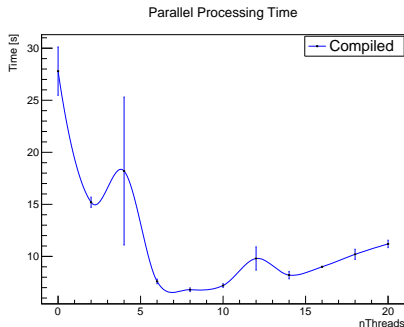
File Sequence Details:

- Run2018_100.root
 - contains 193388 events
 - contains 1547322 unique conversion tracks
- Run2018_110.root
 - contains 140502 events
 - contains 950444 unique conversion tracks
- Run2018_120.root
 - contains 99085 events
 - contains 602330
- Run2018_130.root
 - contains 62289 events
 - contains 347156 unique conversion tracks
- Run2018_141.root
 - contains 218339 events
 - contains 1958272 unique conversion tracks
- Run2018_155.root
 - contains 269483 events
 - contains 2968340 unique conversion tracks
- Run2018_166.root
 - contains 172409 events
 - contains 1299172 unique conversion tracks
- Run2018_176.root
 - contains 127117 events
 - contains 832024 unique conversion tracks
- Run2018_193.root
 - contains 294125 events
 - contains 3685896 unique conversion tracks

First Cluster Test Results:

(Using full 9 file dataset)

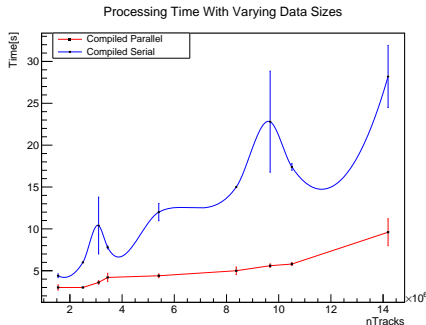
- Time is Mean time $\pm s/\sqrt{N}$ over 5 trials
- data point at nThreads = 0 is the basic sequential program
- high threadcount again not optimal, chokes up the system
- Classic sequential approaches using MakeClass (Compiled and Interpreted)–
 - 265.4 ± 9.2 [s] (Interpreted)
 - 171.0 ± 10.6 [s] (Compiled)
- ROOT6 multithreading is very good!



Test 2: UNL multi-threading

Second Cluster Test: The impact of data size on performance

- Time is Mean time $\pm s/\sqrt{N}$ over 5 trials
- nTracks is the accumulated number of tracks looped over for a given number of files that are run in a preserved order
- Parallel is run with the previously optimal number of threads (8)
- Parallelization is faster and more consistent at any size of dataset
- As data gets larger the parallel benefit becomes more significant



Test 3: UNL multi-processing

Two “new” tools:

Parallel ROOT Facility (PROOF) and **Selectors**

PROOF and Selector are designed to be used together

PROOF is a toolset used to manage/delegate processes between nodes on a cluster

Selector is the ROOT6 version of a MakeClass i.e.

TTreeReaderValues/Arrays

Using them is easy:

- write a simple analysis in Selector
- load all input files into a TChain
- create a PROOF instance
- `tchain.Process(“selector.C”)`

Test 3: UNL multi-processing

Test 3: compare multiprocessing and multithreading at UNL
use the same 9 files with the same task
also try a large 918 file (Single Muon) test case with the same task

the large dataset contains 146,607,893 events

Method	numFiles	nTrials	Time (s)
multiThread	9	5	6.8 ± 0.2
multiProcess	9	3	92.3 ± 0.9
multiThread	918	1	26m 4s
multiProcess	918	1	126m 56s

Multithreading is still significantly faster!!

Process resource and management overhead is high

Possible to combine Multiprocessing + Multithreading in ROOT??

In principle yes..

Application to SUSY Analysis

New tool available at github : [githublink.com](https://github.com)

Basic Analysis template – good for any Tree reading + Histogram producing analysis

Modular tools assembled to automatically parse Trees and perform multithreading → just make histograms!

How it works:

- uses a Python script for input interface
 - just feed in a list of signal or background rootfiles with the treename
- uses a TSelector to access all possible TTreeReaderValue branches in analysis
 - easily automatically generated with just `tree.MakeSelector("myselector")`
- `histset.c` is the meat of the analysis, contains threaded histograms and physics analysis code
 - this is the only place where code needs to be modified
- the custom class `ParTreeProcessing.c` manages the selector/multithreading/histograms

Application to SUSY Analysis

Test 4: Multithreaded SUSY Analysis Running over all large backgrounds and two signal trees

Background: (all HT binned)

DYJetsToLL_HT.list

QCD_HT.list

TTJets_HT.list

WJetsToLNu_HT.list

ZJetsToNuNu_HT.list

Signal:

SMS_TChiWZ_ZToLL.list

SMS_300_75 and SMS_175_135

SUSY Results

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
941908	jangulan	20	0	2438m	2.0g	71m	R	794.9	2.9	5:18.92	compliedthreads
936317	crogan	20	0	754m	213m	48m	R	99.6	0.3	7:48.96	MakeReducedNtup
928604	crogan	20	0	754m	193m	6248	R	99.0	0.3	15:39.93	MakeReducedNtup
927931	crogan	20	0	754m	193m	4472	R	81.7	0.3	15:56.71	MakeReducedNtup
913844	erichjs	20	0	1219m	744m	91m	S	45.9	1.0	19:42.86	python
913694	erichjs	20	0	1223m	748m	91m	S	45.6	1.0	19:12.97	python
1813	erichjs	20	0	6220m	1.0g	3788	S	14.0	2.7	102:7.00	python


```

root [1] _fileio->ls():
File**          susyHists.root
TFile*          susyHists.root
KEY: TH1D       DVMETHist;1      object title
KEY: TH1D       DVCat0NjetSHist;1 object title
KEY: TH1D       DVCat1NjetSHist;1 object title
KEY: TH2D       DVCat0PtcnPtisrDphiCMIHist;1 object title
KEY: TH2D       DVCat1PtcnPtisrDphiCMIHist;1 object title
KEY: TH1D       WJETHist;1      object title
KEY: TH1D       WJETcat0NjetSHist;1 object title
KEY: TH1D       WJETcat1NjetSHist;1 object title
KEY: TH2D       WJETcat0PtcnPtisrDphiCMIHist;1 object title
KEY: TH2D       WJETcat1PtcnPtisrDphiCMIHist;1 object title
KEY: TH1D       QCDMETHist;1    object title
KEY: TH1D       QCDCat0NjetSHist;1 object title
KEY: TH1D       QCDCat1NjetSHist;1 object title
KEY: TH2D       QCDCat0PtcnPtisrDphiCMIHist;1 object title
KEY: TH2D       QCDCat1PtcnPtisrDphiCMIHist;1 object title
KEY: TH1D       TTJETHist;1    object title
KEY: TH1D       TTJcat0NjetSHist;1 object title
KEY: TH1D       TTJcat1NjetSHist;1 object title
KEY: TH2D       TTJcat0PtcnPtisrDphiCMIHist;1 object title
KEY: TH2D       TTJcat1PtcnPtisrDphiCMIHist;1 object title
KEY: TH1D       ZNUNUMETHist;1 object title
KEY: TH1D       ZNUNUCat0NjetSHist;1 object title
KEY: TH1D       ZNUNUCat1NjetSHist;1 object title
KEY: TH2D       ZNUNUCat0PtcnPtisrDphiCMIHist;1 object title
KEY: TH2D       ZNUNUCat1PtcnPtisrDphiCMIHist;1 object title
KEY: TH1D       WZ_300_75METHist;1 object title
KEY: TH1D       WZ_300_75cat0NjetSHist;1 object title
KEY: TH1D       WZ_300_75cat1NjetSHist;1 object title
KEY: TH2D       WZ_300_75cat0PtcnPtisrDphiCMIHist;1 object title
KEY: TH2D       WZ_300_75cat1PtcnPtisrDphiCMIHist;1 object title
KEY: TH1D       WZ_175_135METHist;1 object title
KEY: TH1D       WZ_175_135cat0NjetSHist;1 object title
KEY: TH1D       WZ_175_135cat1NjetSHist;1 object title
KEY: TH2D       WZ_175_135cat0PtcnPtisrDphiCMIHist;1 object title
KEY: TH2D       WZ_175_135cat1PtcnPtisrDphiCMIHist;1 object title


```

plots produced: MET, Njet, S, PTCM/PTISR vs dPhiCMI

$$Z \rightarrow \nu\nu$$



$$Z \rightarrow \ell\ell$$



Final Runtime: Real time= 234.2 ± 10.2 s ≈ (4m)

System time: 61.7 ± 1.6 s

Conclusion

Multithreading is very very fast! – we need to use this –

My analysis template can also be adapted to other tasks i.e. making reduced Ntuples

Next up: Vectorization, Scientific Python - uproot etc.